

2d-transform

A program by Way Yan

Introduction

2d-transform was developed while I was learning multivariable calculus. It is used to visualise functions f from \mathbb{R}^2 to \mathbb{R}^2 . It does so by applying f to the gridlines in the domain, then drawing the result as a pdf file. I use pdf instead of other file formats like png and svg, because pdfs can be scaled without losing quality unlike pngs, and they are natively supported by **pdftex** unlike svgs.

Usage

There should be two shell scripts in the same directory as this file, **compile.sh** and **build.sh**. **build.sh** generates the **L^AT_EX** documentation and C source from **2d-transform.nw** and additionally compiles the C source into the executable **2d-transform**. **compile.sh** only does the second step.

Initially, **2d-transform** will generate an output pdf with a set of default parameters. To run **2d-transform** with different parameters, the user first has to edit the *parameters* section in **2d-transform.c**. Then, the user should run **compile.sh** to generate a new executable. This executable, when run, produces a pdf file based on the parameters specified earlier.

The main reason for requiring the parameters to be specified in the source code is to make my life as a programmer easier. If the function f were to be specified instead as a command-line argument to the executable, I would have to add code to parse the definition of f as a mathematical expression. Though it is certainly not impossible to implement, I find it an unwanted distraction.

The following code outlines the top-level structure of **2d-transform**:

```
<2d-transform.c>≡
  <headers>
  <parameters>
  <internal state>
  <cartesian-to-cairo-coords>
  int main() {
    <setup>
    <draw>
    <cleanup>
  }
```

Headers

`2d-transform` has very few dependencies, hence this section is small. `stdio` consists of standard input/output functions (mainly `printf` for debugging), `math` allows the user to input functions like `sin` in $\langle parameters \rangle$, and `cairo` is the graphics library responsible for drawing. I use `cairo` primarily because it supports pdf output, unlike other libraries.

```
 $\langle headers \rangle \equiv$   
#include <stdio.h>  
#include <math.h>  
#include <cairo.h>
```

Parameters

These are the things that the user has to specify before compiling `2d-transform.c` and running the executable to obtain the pdf.

Of course, the main thing we want to specify is the actual function to draw. These are provided through the parameters `f1` and `f2`, which represent the x and y components of f . If the source code is not modified, then the default function is the identity.

```
 $\langle parameters \rangle \equiv$   
#define f1(x,y) sin(x)  
#define f2(x,y) y+sin(x)  
 $\langle gridline parameters \rangle$   
 $\langle test square parameters \rangle$   
 $\langle misc parameters \rangle$ 
```

Next are the parameters for the drawing of gridlines. Essentially, the domain is defined to be

$$[\text{dom_xmin}, \text{dom_xmax}] \times [\text{dom_ymin}, \text{dom_ymax}].$$

Only the portions of the gridlines lying in the domain are considered for the drawing. Similarly, the codomain is defined to be

$$[\text{codom_xmin}, \text{codom_xmax}] \times [\text{codom_ymin}, \text{codom_ymax}].$$

If you imagine the codomain to be a rectangle in \mathbb{R}^2 containing the transformed gridlines from the domain, then the pdf output will look like that rectangle, but scaled to the dimension specified by `pdfsize`. Also, `gridlinegap` specifies the distance between consecutive horizontal and vertical gridlines in the domain.

Now, each gridline L has infinitely many points—hence one cannot hope to fully determine $f(L)$ by computing $f(x)$ for every single point $x \in L$. Rather, in the case where L is a vertical line, the program will go from top to bottom, sampling a finite number of uniformly spaced points x_i along L (see *draw one vertical gridline* for more details). The spacing between the points x_i is governed by `samplegap`. The smaller `samplegap` is, the more accurate the output will be. Both `gridlinegap` and `samplegap` have no units, because they represent lengths in the coordinate system of the domain.

```
<gridline parameters>≡
#define dom_xmin -5
#define dom_xmax 5
#define dom_ymin -5
#define dom_ymax 5
#define codom_xmin -5
#define codom_xmax 5
#define codom_ymin -5
#define codom_ymax 5
#define gridlinegap 0.5
#define samplegap 0.1
```

`2d-transform` also has a feature to visualise the image of a specified square under f . It is enabled by setting `tseenable` to 1 and disabled by setting it to 0. The difference between this and `gridlines` is that the area of the square is filled in with points, and also the sides and corners of the square are drawn in such a way that they can be distinguished. Hence this feature reveals some features of f that are otherwise hard to tell, such as orientation. It is also helpful when there are many gridlines mapping to the same region in the domain.

The square, as a subset of the domain, is defined as

$$[\text{ts_xmin}, \text{ts_xmax}] \times [\text{ts_ymin}, \text{ts_ymax}].$$

Also, `tsgap` specifies the gap to leave between points when filling up the area of the square. As with `gridlinegap` and `samplegap`, the parameter `tsgap` has no units because it is relative to the coordinate system of the codomain.

```

<test square parameters>≡
#define tseenable 1
#define ts_xmin -2
#define ts_xmax 2
#define ts_ymin -2
#define ts_ymax 2
#define tsgap 0.4

```

The original and transformed test square is shown in Figures 1 and 2. Note that as one goes from left to the right, the size of the circles transitions from `ts_startsize` to `ts_endsize`. Likewise as one goes from bottom to top, the colour transitions from `(ts_startr, ts_startg, ts_startb)` to `(ts_endr, ts_endg, ts_endb)`.




Figure 1: How the square looks like before transformation.




Figure 2: How the square looks like after the transformation $(\sin x, y + \sin x)$.

```
<test square parameters>+=
#define ts_startsize 1
#define ts_endsize 3
#define ts_startr 255
#define ts_startg 0
#define ts_startb 0
```

```

#define ts_endr 0
#define ts_endg 0
#define ts_endb 255

```

Lastly we have miscellaneous parameters. `pdfsize` is width of the pdf, in points. (For reference, an A4-sized pdf is 595 points wide and 842 points high.) Actually, `pdfsize` is also the height of the pdf, because I make the width and height equal for simplicity. `outputfile` is the name of the pdf file to save the drawing to.

```

<misc parameters>≡
#define pdfsize 100
#define outputfile "out.pdf"

```

Internal state

These are variables that are used internally, and should be not be touched by the user. Firstly, I provide the function definition of `cairo_pdf_surface_create`; otherwise, the compiler throws an `Implicit function declaration` error (I'm not fully sure why). Then, `surface` and `cr` is used by `cairo` for drawing.

Lastly, `pair` is effectively a register variable that stores the output after calling `cartesian_to_cairo_coords`. It is implemented this way, because C does not allow `cartesian_to_cairo_coords` to return an array directly. It can return a pointer to an array, but I find it completely unnecessary to deal with pointers for such a small array. Hence I resorted to the simplest solution.

```

<internal state>≡
cairo_surface_t *cairo_pdf_surface_create(const char *filename,
                                           double width_in_points, double height_in_points);
cairo_surface_t *surface;
cairo_t *cr;
double pair[2];

```

Changing coordinate systems

The coordinate system that `cairo` uses is different from the Cartesian coordinate system: the top-left corner is $(0,0)$ while the bottom-right corner is $(\text{pdfsize}, \text{pdfsize})$. In Cartesian coordinates, $(0,0)$ would be at the center.

The function takes in the Cartesian coordinates (x,y) as input. It computes the values $t_x, t_y \in [0,1]$ such that

$$\begin{aligned}x &= (1 - t_x) * \text{codom_xmin} + t_x * \text{codom_xmax} \quad \text{and} \\y &= (1 - t_y) * \text{codom_ymax} + t_y * \text{codom_ymin}.\end{aligned}$$

Essentially, t_x measures how far right (x,y) should be in the pdf file from the left; t_y measures how far down it should be from the top. Note that going up in a pdf file corresponds to decreasing the y coordinate. Then, we simply scale t_x and t_y by `pdfsize`.

```
<cartesian-to-cairo-coords>≡
void cartesian_to_cairo_coords(double x, double y) {
    double tx = (x-codom_xmin)/(codom_xmax-codom_xmin);
    double ty = (y-codom_ymax)/(codom_ymin-codom_ymax);
    pair[0] = tx * pdfsize;
    pair[1] = ty * pdfsize;
}
```

Setup and cleanup

Nothing much to say really.

```
<setup>≡
surface = cairo_pdf_surface_create(outputfile, pdfsize, pdfsize);
cr = cairo_create(surface);
cairo_set_source_rgb(cr, 0, 0, 0);
cairo_set_line_width(cr, 0.1);

<cleanup>≡
cairo_destroy(cr);
cairo_surface_destroy(surface);
return 0;
```

Drawing

This is the real meat of the program.

```
<draw>≡  
double xcur, ycur;  
  <draw transformed vertical gridlines>  
  <draw transformed horizontal gridlines>  
  if (tsenable) {  
    <draw test square>  
  }
```

```
<draw transformed vertical gridlines>≡  
  <find leftmost vertical gridline>  
  <iterate through vertical gridlines>
```

Besides being equally spaced according to `gridlinegap`, I require that one of the vertical gridlines be the y -axis. Doing so fixes all the other vertical gridlines in \mathbb{R}^2 and hence D .

With this, the leftmost vertical gridline is the n th one from the y -axis. If $n > 0$, then it lies to the right of the y -axis. If $n < 0$ then it lies to the left.

```
<find leftmost vertical gridline>≡  
int n = dom_xmin / gridlinegap;  
xcur = n * gridlinegap;
```

The vertical gridlines which intersect D are drawn.

```
<iterate through vertical gridlines>≡  
while (xcur <= dom_xmax) {  
  <draw one vertical gridline>  
  cairo_stroke(cr);  
  xcur += gridlinegap;  
}
```


The point (xcur, ycur) iterates through the sampled points x_i (see *parameters*). Lines are drawn connecting each pair $(f(x_i), f(x_{i+1}))$, the idea being that the resulting collection of line segments closely resembles $f(L)$.

```

<draw one vertical gridline>≡
    ycur = dom_ymin;
    cartesian_to_cairo_coords(f1(xcur, ycur), f2(xcur, ycur));
    cairo_move_to(cr, pair[0], pair[1]);
    while (ycur <= dom_ymax) {
        cartesian_to_cairo_coords(f1(xcur, ycur), f2(xcur, ycur));
        cairo_line_to(cr, pair[0], pair[1]);
        ycur += samplegap;
    }

```

The code for drawing horizontal gridlines is essentially the same except the characters x and y are swapped. Hence I will not divide it further.

```

<draw transformed horizontal gridlines>≡
    n = dom_ymin / gridlinegap;
    ycur = n * gridlinegap;
    while (ycur <= dom_ymax) {
        xcur = dom_xmin;
        cartesian_to_cairo_coords(f1(xcur, ycur), f2(xcur, ycur));
        cairo_move_to(cr, pair[0], pair[1]);
        while (xcur <= dom_xmax) {
            cartesian_to_cairo_coords(f1(xcur, ycur), f2(xcur, ycur));
            cairo_line_to(cr, pair[0], pair[1]);
            xcur += samplegap;
        }
        cairo_stroke(cr);
        ycur += gridlinegap;
    }

```

Next, the test square is drawn.

```
<draw test square>≡
double curx, cury;
double tx, ty, r, g, b, size;
curx = ts_xmin;
while (curx <= ts_xmax) {
    cury = ts_ymin;
    while (cury <= ts_ymax) {
        <draw circle>
        cury += tsgap;
    }
    curx += tsgap;
}

<draw circle>≡
tx = (curx-ts_xmin)/(ts_xmax-ts_xmin);
ty = (cury-ts_ymin)/(ts_ymax-ts_ymin);
r = ((1-ty)*ts_startr + ty*ts_endr) / 255;
g = ((1-ty)*ts_startg + ty*ts_endg) / 255;
b = ((1-ty)*ts_startb + ty*ts_endb) / 255;
size = (1-tx)*ts_startsize + tx*ts_endsize;
cairo_set_source_rgba(cr, r, g, b, 0.5);
cartesian_to_cairo_coords(f1(curx,cury), f2(curx, cury));
cairo_arc(cr, pair[0], pair[1], size, 0, 2*M_PI);
cairo_fill(cr);
```